

Story-by-Story Board Review

Every one of the 73 user stories on the Azure DevOps board, reviewed against the code, the database, the pipelines and the test suites. What exists, what is tested, and what is still needed.

REVIEWER	DATE	BOARD	DAYS TO SUBMIT
Shanthi Subramanian · BA / QA	22 September 2026	91 items · 18 epics · 73 stories	8

73	100%	31	28	14
STORIES REVIEWED	HAVE CRITERIA	WELL COVERED	PARTIAL	LITTLE OR NONE

Summary

The headline findings before the detail.

THE BACKLOG ITSELF IS IN GOOD SHAPE

All 73 stories have acceptance criteria — not one is empty. They average five criteria each, they are written in Given/When/Then, and the platform stories (Epics 15–18) are unusually specific for infrastructure work. This is a better-specified backlog than most delivery teams have, and it is the main reason a review like this is possible at all.

Coverage at a glance

GROUP	STORIES	WELL COVERED	PARTIAL	LITTLE / NONE
Product — Epics 1–10	24	11	8	5
QA-raised — Epics 11–13	6	0	2	4
Mock systems — Epic 14	6	4	2	0
Platform — Epics 15–18	36	16	15	5
Documentation	1	0	1	0

The five things that matter most

- US-9.2 (Guarantee immutability)** — the flagship governance claim, implemented by a database trigger, and **tested by nothing at all**. No HTTP endpoint can reach it.
- US-12.1 (Apply schema to a deployed environment)** — still New, still unowned, and now blocking a live deployment attempt.

3. **US-11.1 (Role-based permissions)** — not implemented. Five API tests exist and all record the gap.
4. **US-17.9 / 17.10 (Build and deploy the container images)** — the pipeline files are **zero bytes** and are not registered in Azure DevOps.
5. **US-12.2 and US-12.3 are marked but the work is not complete** — closed as superseded by Epics 17 and 18, which is reasonable, but reads as delivered.

Review method

What was consulted, and how each story was judged. Stated so the conclusions can be checked rather than taken on trust.

Inputs

SOURCE	USED FOR
Azure DevOps REST API wit/workitems, /revisions, /comments	All 91 items, their criteria, revision history and comments — so "who changed what, when" is evidence rather than recollection
Git repository main · release/1.00 · 6 feature branches	Whether each story is actually implemented, and on which branch
Database schema schema/*.sql	Allowed values and numeric bounds read from CHECK constraints and triggers — the contract the system truly enforces, not documentation
Test suites xUnit · Playwright · evals	Executed locally; 62 unit, 18 Playwright, 12 eval cases mapped back to stories
Pipeline definitions + build history build/definitions, build/builds	What actually runs automatically, versus what exists as a file in the repo
The running application	Manual exploration of all six screens and the seven workspace tabs

How each story was rated

Two independent axes. A story can be fully built and untested, or well specified and unbuilt — conflating them hides exactly the problems worth finding.

AXIS	VALUES	JUDGED BY
Built	YES PARTIAL NO	Reading the implementing code, schema or pipeline — never the story's state field
Tested	●●● strong · ●●○ partial · ●○○ little or none	An executed test that names the story or covers its criteria

●●● means covered at *more than one layer*, including one that cannot be bypassed. ●●○ means one layer only, or a known gap inside the story. ●○○ means no automated coverage that would catch a regression.

ONE DELIBERATE RULE

The story's state field was ignored when judging "Built". Two stories are marked Done whose work is incomplete, and many marked New are fully implemented — the board's state reflects workflow bookkeeping, not delivery. Reading the code is the only reliable method.

Epics 1–5 · Intake through AI draft

The core product flow. Largely built; the AI touchpoints carry the coverage risk.

Epic 1 Change Request Intake

ID	STORY	BUILT	TESTED	WHAT'S THERE	WHAT'S NEEDED
2	US-1.1 Submit a change request	YES	●●●	Six change types, immutable <code>CR-2026-nnnnn</code> , type-specific JSON fields. Covered by UI journey 1, unit test on ingestion trigger	Nothing urgent
3	US-1.2 Attach supporting documents	YES	●●○	PDF/DOCX/XLSX enforced by DB constraint; 6 unit tests on content type and upload ordering	No test that replacing a document versions it rather than overwriting — a stated criterion
4	US-1.3 Track my request status	PARTIAL	●●○	Status and days-elapsed list; covered by UI journey 1	Notifications are not implemented (criterion 2). Error-handling criteria added for <code>DEF-003</code> are untested and currently failing

Epic 2 Risk Categorization & Framework Mapping

ID	STORY	BUILT	TESTED	WHAT'S THERE	WHAT'S NEEDED
6	US-2.1 Auto-map to risk categories <code>AI</code>	YES	●●○	FFIEC-grounded proposal; the client discards any category outside the allowed list — a structural guard. 5 unit tests, 5 eval cases. Francis's new criterion on grounding against the snapshot is supported by the code	4 of 5 eval cases blocked on a model endpoint. <code>DEF-006</code> criterion untested and failing
7	US-2.2 Override proposed categories	YES	●○○	Override endpoint writes an audit row in the same transaction	No test that a blank reason is refused — the story's central criterion. One API test would close it

Epic 3 AI-Assisted Policy Research

ID	STORY	BUILT	TESTED	WHAT'S THERE	WHAT'S NEEDED
9	US-3.1 Surface relevant policy <code>AI</code>	YES	●○○	Postgres full-text search — <i>deliberately not an LLM</i> , and a good "when not to use AI" decision	No coverage at all. Needs: ranked results returned, and "no matching policy found" flagged explicitly rather than left blank
10	US-3.2 Confirm or reject surfaced policy	YES	●●○	Reliance decisions recorded; finalization blocked without one — proven by unit test UT-AS-03	No test of the <i>recording</i> itself, only of the downstream gate

Epic 4 AI-Assisted Document Extraction

ID	STORY	BUILT	TESTED	WHAT'S THERE	WHAT'S NEEDED
12	US-4.1 Extract structured data AI	YES	●●○	Each field carries a source excerpt and a needsReview flag; 3 unit tests, 4 eval cases; real PDF/DOCX text extraction	3 of 4 eval cases blocked on the model. No test that the source page/section pointer is populated
13	US-4.2 Correct extracted data	YES	●○○	Correction endpoint retains the original AI value alongside the correction	No test that both values are retained — this is the "never a silent overwrite" rule and it is unproven

Epic 5 AI-Drafted Risk Assessment

ID	STORY	BUILT	TESTED	WHAT'S THERE	WHAT'S NEEDED
15	US-5.1 Generate a draft assessment AI	YES	●●●	Per-category narrative; unsupported claims returned in a separate field, never folded into the text; labelled "AI-drafted" and blocked from committee until reviewed (UT-AS-01)	Narrative <i>quality</i> is a heuristic proxy only — flagged as manual scenario MV-01
16	US-5.2 Regenerate a section with feedback	YES	●●○	Feedback reaches the prompt (unit test UT-AI-11); eval case nr-03	No test that other sections are untouched by a regeneration, or that the prior version stays retrievable

Epics 6–10 · Review, scoring, committee, audit, config

Where the governance guarantees live — and where the single largest gap sits.

Epic 6 Analyst Review, Edit & Override

ID	STORY	BUILT	TESTED	WHAT'S THERE	WHAT'S NEEDED
18	US-6.1 Edit any AI-generated output	YES	●○○	Every service method takes a mandatory reason; stored functions write the audit row in the same transaction	No test that a blank reason is refused. This is the cross-cutting override rule for the whole product and it is unproven at any layer
19	US-6.2 View full edit/override history	YES	●●○	Audit tab shows the chain; UI journey 2 asserts the intake event appears attributed to a named human	No test that a full chain shows original, edited value, actor, timestamp and reason together
20	US-6.3 Finalize the assessment	YES	●●●	Best-covered story on the board: 12 unit tests, 2 UI tests, verified by mutation testing	DEF-004 — an assessment with no categories still finalizes. Criteria added 17 Sep; fix outstanding

Epic 7 Risk Scoring Engine

ID	STORY	BUILT	TESTED	WHAT'S THERE	WHAT'S NEEDED
22	US-7.1 Risk score with breakdown	PARTIAL	●●○	Deterministic formula; residual > 0 guaranteed by arithmetic and enforced in three places. 2 unit tests	DEF-005 — the <code>control</code> table is never seeded, so "trace which controls were credited" cannot be satisfied . Needs a control catalogue (BA input) plus the 9-row scoring truth table
23	US-7.2 Override a calculated score	YES	●●●	3 unit tests: non-positive residual refused, blank reason refused, valid override accepted. DB constraint as backstop	No test that an overridden score is visually distinguishable from a calculated one

Epic 8 Committee Review & Voting

ID	STORY	BUILT	TESTED	WHAT'S THERE	WHAT'S NEEDED
25	US-8.1 Route a finalized assessment	YES	●●●	UI journeys 2 and 3 cover routing and queue arrival	No test that the committee view is genuinely read-only
26	US-8.2 Cast a committee vote	YES	●●●	14 unit tests on quorum and the conservative resolution rule; UI journey 3 asserts the vote is named	DEF-002 — any caller can vote as anyone. Criterion added 17 Sep; API test records the gap
27	US-8.3 Record the final decision	YES	●●●	All four resolutions covered by unit tests; conditions text merged from every vote, nothing dropped	No test that the Product Owner is notified , or that conditions stay attached to the request record

Epic 9 Immutable Audit Trail

ID	STORY	BUILT	TESTED	WHAT'S THERE	WHAT'S NEEDED
29	US-9.1 Reconstruct full history	YES	●●○	Audit endpoint returns the chain; UI journey 2 checks one event	No test that every event type appears in order with before/after state
30	US-9.2 Guarantee immutability	YES	●○○	A database trigger rejects <code>UPDATE</code> and <code>DELETE</code> on <code>audit_event</code> outright — enforced at the data layer exactly as the story requires	Tested by nothing. No endpoint attempts to alter an audit row, so no API or UI test can reach it. This is the single most important gap on the board — needs the Testcontainers suite (DT-01... 03)
38	US-9.3 Export an examiner-ready trail QA-RAISED	PARTIAL	●○○	Being built now on <code>fix/azure-deploy-and-audit-export</code> — new <code>AuditExportService</code> , PDF export, controller endpoint and UI wiring	No tests yet. Examiner-readiness is manual scenario MV-02 — no assertion can judge "suitable for handing to an examiner"

Epic 10 Platform Configuration

ID	STORY	BUILT	TESTED	WHAT'S THERE	WHAT'S NEEDED
32	US-10.1 Configure scoring parameters	YES	●●●	6 unit tests on the mitigation-factor bounds plus an API test; versioned config so in-flight assessments keep their original values	No test that an in-flight assessment is genuinely unaffected by a config change
33	US-10.2 Configure workflow rules	YES	●○○	Workflow rules stored as JSON with a mandatory reason; the committee quorum is read from here	No coverage. Needs: change is audited, reason mandatory, in-flight requests follow the rules active when submitted

Epics 11-14 · QA-raised and mock systems

Six stories QA raised on 17 September, and the mock-systems epic added by Dev.

Epics 11-13 Access Control · Deployment & Operations · Non-Functional

ID	STORY	BUILT	TESTED	WHAT'S THERE	WHAT'S NEEDED
37	US-11.1 Enforce role-based permissions	NO	●●○	Roles are seeded and the frontend filters navigation by role. 5 API tests exist and all record the gap	The API enforces nothing. The acting user is read from the request body. Needs the authorization layer DEF-001 DEF-002
39	US-12.1 Apply schema to a deployed env	NO	●○○	Nothing. The DB folder is not in either image, the app performs no bootstrap, no pipeline applies it	First-deploy blocker. PostgreSQL is now provisioned in Azure and empty. Still unowned — open question Q6 DEF-007
40	US-12.2 Build and publish images DONE	PARTIAL	●○○	Containers now run non-root (fixed 21 Sep); <code>.dockerignore</code> keeps secrets out of layers	No pipeline builds, scans or pushes an image. Closed as "covered by Epic 17" — see the state-accuracy note below
41	US-12.3 Observe a running environment DONE	PARTIAL	●○○	Both services register OpenTelemetry and activate on a connection string; the Log Analytics module now exists	<code>log_analytics_workspace_id</code> is still commented out in <code>main.cae.tf</code> . Closed as "covered by Epic 18"
42	US-13.1 Retain records for the supervisory period	NO	●○○	Nothing — no retention or backup policy defined	A stated retention period (5 years minimum) and point-in-time restore on the Flexible Server
43	US-13.2 Keep model calls inside the tenant	PARTIAL	●○○	Azure AI Foundry is provisioned and the provider is swappable with no prompt changes	The default provider is still the non-Azure path. Needs a config assertion that a deployed environment cannot select it

STATE ACCURACY — WORTH RAISING WITH SURAJ

US-12.2 and US-12.3 were moved to **Done** on 21 September with the comments *"Covered as part of epic 17"* and *"Covered as part of Epic 18"*. Closing them as superseded is entirely reasonable — Epics 17 and 18 do cover that ground.

But **Done** means **delivered**, and neither is: the image pipeline files are zero bytes and unregistered in Azure DevOps, and Log Analytics is still commented out. In a Scrum process the state for superseded work is **Removed**, with a link to the superseding item. As it stands the burndown is inflated and anyone scanning the board reads image publishing and observability as complete.

Epic 14 Mock External Systems & Data Ingestion

ID	STORY	BUILT	TESTED	WHAT'S NEEDED
45	US-14.1 Stand up the three mock systems	YES	●●●	Separate service, own schema, own port; API smoke test confirms 6 customers / 5 products / 4 vendors
46	US-14.2 Fields the framework scores against	YES	●●○	No test that each mock record carries the fields the FFIEC categories actually need
47	US-14.3 Look up an existing record at intake	YES	●●●	Covered by the three <code>/api/DataIngestion/Mock*</code> endpoints and smoke tests
48	US-14.4 Ingest and snapshot at intake	YES	●●●	5 unit tests including partial resolution and outage tolerance. Snapshot immutability untested — belongs in the Data layer
49	US-14.5 Push the decision back to the source	YES	●●●	4 unit tests; a Mock Systems outage correctly does not fail the vote
50	US-14.6 Generate mock data reproducibly	YES	●●○	Seed exists and is documented in <code>ai/data-generation/</code> ; no test that a re-seed is reproducible

Epics 15–16 · Terraform modules & dev environment

23 stories. Well specified, largely built, and covered by the IaC pipeline rather than by QA test suites.

HOW THESE ARE TESTED

Infrastructure is not tested by QA test suites — it is tested by the **IaC PR pipeline**, which now hard-gates on `terraform validate`, TFLint error-severity, Trivy HIGH/CRITICAL and a successful `terraform plan`. Checkov remains soft-fail by a documented decision with every suppression justified. That is a reasonable model and it runs on every infrastructure PR, which is more than the application layers currently get.

Epic 15 Terraform Modules · 10 stories

ID	MODULE	BUILT	TESTED	NOTES
53	US-15.1 Resource Group	YES	●●●	Six mandatory tag validations — a real governance rule, enforced in the module
54	US-15.2 Virtual Network	YES	●●●	Subnets, NSGs, route tables, peering; four worked examples
55	US-15.3 Private DNS Zone	YES	●●○	Built and wired
56	US-15.4 Key Vault	YES	●●○	8 criteria — the most security-relevant module; role assignments wired
57	US-15.5 Storage Account	YES	●●○	9 criteria
58	US-15.6 PostgreSQL Flexible Server	YES	●●○	Random password generated; no backup/retention criterion — links to US-13.1
59	US-15.7 Log Analytics Workspace	YES	●●○	Module exists; consumption is the gap, not the module
60	US-15.8 Container Registry	YES	●●○	<code>admin_enabled = false</code> ; six Checkov suppressions, each justified in writing
61	US-15.9 Container Apps Environment	YES	●●○	Supports VNet integration, internal load balancer and DNS linking
62	US-15.10 Container Apps	YES	●●○	9 criteria; env vars and identity wired

Epic 16 Dev Environment · 13 stories

ID	STORY	BUILT	NOTES
64–69	US-16.1...16.6 State, RG, VNet, DNS, Key Vault, Storage	YES	All wired into <code>iac/environments/dev</code> . The VNet is now genuinely consumed — the CAE takes the delegated subnet DEF-008 CLOSED
70	US-16.7 PostgreSQL server and database	YES	Provisions the server; nothing populates the schema. This is where US-12.1 bites
71	US-16.8 Log Analytics + Application Insights	PARTIAL	Workspace module exists; not connected to the Container Apps Environment
72–74	US-16.9...16.11 Registry, managed identity, CAE	YES	Identity-based image pull — no stored registry credential
75–76	US-16.12/16.13 Workbench and Mock API container apps	YES	Renamed to <code>gh-hrg-workbench</code> / <code>gh-hrg-mockapi</code> ; env vars wired

Epics 17–18 • CI/CD and observability

13 stories. Most are genuinely delivered; three are not, and one of those blocks deployment.

Epic 17 DevOps CI/CD Pipelines • 11 stories

ID	STORY	BUILT	EVIDENCE
78	US-17.1 Azure DevOps prerequisites	YES	Service connections and variable groups in use by four live pipelines
80	US-17.2 Infra PR pipeline: fmt, validate, lint, scans	YES	Running; hard gates added 21 Sep
81	US-17.3 Infra PR: plan without stored credentials	YES	Workload Identity Federation — no client secret anywhere
82	US-17.4 AI review for infra PRs with verdict gate	YES	Running; verdict derived deterministically, PR author now @mentioned
83	US-17.5 Infra deploy pipeline: apply on merge	YES	<code>gh-infra-dev-deploy</code> succeeded 21 Sep
84	US-17.6 Workbench PR: build, tests, audit, SAST	YES	<code>gh-workbench-dev-pr-review</code> succeeded 21 Sep. Gates unit tests only — no API, Data or UI stage
85	US-17.7 Mock API PR: SQL drift check, build, audit	YES	Succeeded 22 Sep. The <code>deploy_all.sql</code> drift check is a nice touch
86	US-17.8 Shared AI code review for app PRs	YES	Refactored into shared templates and scripts
87	US-17.9 Workbench: build, publish, deploy image	NO	<code>workbench.yml</code> is 0 bytes and not registered as a pipeline
88	US-17.10 Mock API: build, publish, deploy image	NO	<code>mock-api.yml</code> is 0 bytes and not registered
89	US-17.11 Enforce as required checks; document	PARTIAL	<code>.azure-pipelines/README.md</code> exists; branch-protection status not verifiable from here

Epic 18 Observability • 2 stories

ID	STORY	BUILT	WHAT'S NEEDED
91	US-18.1 Send logs and metrics to App Insights	PARTIAL	Application code is instrumented and inert until a connection string is set. <code>log_analytics_workspace_id</code> is one commented line from working
92	US-18.2 SRE agent watching App Insights	NO	Not started. Depends on US-18.1; the most deferrable item on the board with 8 days left

Cross-cutting findings

Patterns that only appear when you read all 73 stories together.

1 • The override rule is the least-tested rule in the product

"A human can override any AI output; must state a reason; both original and override are retained" is a **non-negotiable constraint** in the project brief and appears in US-2.2, US-4.2, US-6.1, US-7.2 and US-10.1. Of those five, **only US-7.2 has a test that a blank reason is refused**.

The implementation looks right everywhere — every service method takes a reason and the stored functions write the audit row in the same transaction. But "looks right" is not "proven", and this is the rule the 15% human-in-the-loop criterion rests on. **Five API tests would close it.**

2 • Data-layer guarantees are systematically untested

Four separate stories depend on something the database enforces: audit immutability (US-9.2), residual risk above zero (US-7.1), snapshot immutability (US-14.4), and config versioning (US-10.1). **None of the four has a test that reaches the database.** One Testcontainers suite closes all of them at once — which is why it has been the top recommendation since 12 September.

3 • The platform epics are better specified than the product epics

Epics 15–18 average **6.6 criteria per story**; Epics 1–10 average **3.2**. The infrastructure stories are also more precise — they name resources, settings and expected outputs. That is unusual and worth keeping; it is why this review could assess them against the Terraform directly.

4 • Notifications appear in three stories and are implemented in none

US-1.3, US-8.1 and US-8.3 each require the Product Owner to be notified. There is no notification mechanism in the codebase. This is a coherent scope decision if made deliberately — but it is currently an *implicit* one, and three stories cannot be marked complete until it is either built or the criteria are removed with a stated reason.

5 • Board state does not reflect delivery

Seventy-one of 73 stories are **New**, including US-6.3 and US-8.2 which are among the best-implemented and best-tested on the board. The two marked **Done** are the two whose work is *not* complete. The state field is currently the least reliable signal on the board — worth a five-minute pass before any demo or review.

What to do next

Eight days. Ordered by what protects the most, for the least effort.

#	ACTION	CLOSES	EFFORT	OWNER
1	Schema deployment to Azure	US-12.1 · unblocks any deploy at all	Medium	DevOps + Dev
2	Testcontainers suite — 24 tests	US-9.2, plus data guarantees in US-7.1, US-10.1, US-14.4	Medium	QA
3	Five override-reason API tests	US-2.2, US-4.2, US-6.1, US-10.2 — the cross-cutting rule	Low	QA
4	Fill <code>workbench.yml</code> / <code>mock-api.yml</code>	US-17.9, US-17.10 · nothing deploys without them	Medium	DevOps
5	Wire API + Data suites into the PR pipeline	Turns 30 written tests into gates	Low	DevOps + QA
6	Fix DEF-002 and DEF-004	US-11.1, US-6.3 · the two High governance defects	Medium	Dev
7	Uncomment Log Analytics on the CAE	US-18.1, US-16.8	Very low	DevOps
8	Seed the control catalogue	US-7.1 explainability DEF-005	Low	Dev + BA
9	Board hygiene: states, and Removed vs Done	Makes the board readable for a reviewer	Very low	All

IF ONLY THREE THINGS GET DONE

Items 1, 2 and 3. Item 1 is the difference between a deployed product and a broken one. Item 2 converts the product's flagship governance claim from an assertion into evidence. Item 3 is a few hours of work that covers the constraint the brief calls non-negotiable across four stories at once.

Decisions still needed from the team

QUESTION	BLOCKS	RAISED
Q3 — does a scoring override need secondary sign-off?	US-7.2, US-6.3 criteria	Original backlog
Q5 — who may approve configuration changes?	US-10.1, US-10.2 criteria	Original backlog
Q6 — who owns database schema deployment?	US-12.1 — <i>the deploy blocker</i>	17 Sep · still unanswered
New — are notifications in scope?	US-1.3, US-8.1, US-8.3	This review

Reviewed 22 September 2026 against the Azure DevOps board (91 items), `origin/main` at `863ead0`, `origin/release/1.00` at `c0705b0`, six feature branches, the database schema, four live pipeline definitions with their build history, and locally executed test suites. "Built" was judged from the implementing code rather than from the story's state field. Where a conclusion could not be verified from here — branch protection settings, for instance — it is marked as such rather than assumed.